

Counting Fruits with State-of-the-Art Multi-View-Stereo Approaches

Bachelor Thesis in Computer Science

presented by

Alexander Link

born on 29.08.2004 in Bamberg


submitted to

Department of Computer Science
Professorship for Cognitive Computer Vision
Friedrich-Alexander-Universität Erlangen-Nürnberg

Academic Supervision: Prof. Dr. Bernhard Egger

Begin date: 01.05.2026

End date: 01.10.2026



Abstract

Contents

1	Introduction	1
2	Related Work	3
2.1	Fruit Counting and Yield Estimation	3
2.2	3D-Reconstruction	4
3	Method	5
3.1	Tense	5
3.2	I, We, Passive voice	5
3.3	Equations	5
3.4	Figures	5
3.5	Tables	6
4	Implementation	7
4.1	Listings	7
5	Results	9
6	Conclusions	11
6.1	Tips	11

CONTENTS

Chapter 1

Introduction

The automatic counting of fruit from images is a key task in precision agriculture, where accurate yield estimates support harvest planning and resource allocation. Reconstructing a fruit tree scene in 3D from a set of RGB images and counting the fruit it contains is appealing because it requires only inexpensive and widely available camera hardware, rather than specialized depth sensors.

citation

Recent Work such as FruitNeRF [10] show that this is possible by training a Neural Radiance Field [11] per scene and counting the fruits in the resulting volumetric representation. This approach is effective but requires a time-consuming scene-specific optimization. In this work, we instead investigate whether a feed-forward multi-view stereo model, VGGT [13], can replace scene-by-scene training. Instead of optimizing a representation of each scene, VGGT predicts a 3D reconstruction in a single forward pass. This promises a significant reduction in runtime, but raises the central question of whether the resulting geometry is sufficiently accurate for reliable counting.

Counting fruits in this scenario is challenging for several reasons: Fruits are often obscured by foliage and each other, the appearance of a scene varies with changing lighting, and the geometric complexity of a tree canopy makes a consistent reconstruction difficult. To reduce occlusion and the risk of missing apples, a scene should ideally be captured from all sides, but the number of views is limited by hardware constraints such as available memory. The cameras are therefore sparsely distributed across a full 360-degree range, and widely spaced views often capture completely different regions of the tree. This makes it difficult to establish correspondences between views

This results in two types of opposing types of errors. Individual apples observed by only a few cameras yield unstable correspondences and are easily split into multiple detections, while dense clusters of apples seen by many cameras may merge into a single detection. A single set of static parameters can therefore rarely resolve both cases optimally at the same time.

Our approach maps semantic information in the form of 2D instance masks, obtained using SAM3 [2], across views via the 3D representation predicted by VGGT. The fruits are then counted by combining a clustering step with a preceding partial instance association that supports the clustering.

Chapter 2

Related Work

2.1 Fruit Counting and Yield Estimation

Approaches to fruit counting can be distinguished based on whether they operate in the image plane or in a three-dimensional representation of the scene. Traditional image-based pipelines are divided into two subtasks: the detection of fruits in the image and their counting across multiple views [5]. While detection used to rely on handcrafted color and shape features, it is now dominated by deep neural networks [12]. For example, YOLO-based detectors achieve high accuracies on single images [1]. However, such methods count fruits on a per-frame basis and lack consistent instance tracking across multiple views, meaning that the same fruit can be counted multiple times. It is precisely this lack of cross-view tracking that is the problem addressed by the 3D-based approach pursued in this work.

This double counting is the central challenge in counting across image sequences. Liu et al. [8] introduced the first complete pipeline that tracks fruits across consecutive images and uses Structure-from-Motion (SfM) to correct double counting, more recent pipelines combine modern detectors such as YOLO with dedicated tracking modules for the same purpose [3]. In a subsequent paper, Liu et al. [9] combined detection with a semantic SfM reconstruction and assigned detections to one another using the Hungarian algorithm, using only a single camera. The association using the Hungarian algorithm is also adopted in the present work; however, there it operates on densely sampled video sequences with short frame intervals, whereas the scenario considered here is based on widely spaced 360-degree views with weak correspondences between the views.

To structurally avoid double counting, later studies shifted the counting entirely into three-dimensional space. Häni et al. [5] reconstructed rows of trees from both sides and merged the counts by combining the 3D models. Gené-Molá et al. [4] combined Mask R-CNN instance segmentation with an SfM point cloud, projecting the 2D masks into 3D space, and filtered out false detections using a support vector machine. This approach is the closest in structure to the present work, as it maps 3D instance masks to views via a 3D representation. However, it relies on an SfM reconstruction, which takes several hundred minutes per scene to complete and is therefore not suitable for a fast run, precisely the limitation that this work seeks to overcome by using a feed-forward model.

FruitNeRF [10] fits into this trend as the approach that fully integrates counting in to a neural scene representation, thereby overcoming the previous separation of reconstruction and segmentation. However, since the NeRF representation used must be optimized per scene, the computational cost remains

high. This limitaion forms the immediate starting point of this work. More recent works address this by replacing per-scene optimization with faster representations: FruitGaussian [6] and CountingFruit [7] use 3D Gaussian splatting to significantly speed up reconstruction, but they remain within the framework of scene-specific optimization.

2.2 3D-Reconstruction

Chapter 3

Method

3.1 Tense

Mind your use of present and past tense. Reports are typically written in present tense, except for the related work and conclusion section, which are usually in past tense. If you describe what past authors did, use past tense. If you describe what an algorithm does, which still works the same way today, use present tense.

3.2 I, We, Passive voice

It is very uncommon to write “I”. Instead, either avoid using personal pronouns by switching to a passive voice, or use “we” instead. When reading the word “we”, you can think of this as “the reader and me”.

3.3 Equations

For equations, the `align` command is recommended, since it allows you to align equations at a specific location, such as at the `=` sign.

$$y = f(x) \tag{3.1}$$

$$\boldsymbol{f} = \boldsymbol{m} \cdot \boldsymbol{a}^\top \tag{3.2}$$

Equations are referenced using `eqref`, for example Equation (3.1) and Equation (3.2). See `macros.tex` for a range of convenient macros, e.g., vector $\boldsymbol{x}$ and matrix $\boldsymbol{M}$ symbols

3.4 Figures

Figure placement is best done with the default `figure` environment. To reference a figure use `figref`, for example Figure 3.1. If possible, always create figures as vector graphics. For example an overview / pipeline graphic of your method should always be a vector graphic. Vector graphics allow zooming into figures without a loss in quality of the figure. Vector graphics can be exported in PDF format from common editing programs, which allows including them in Latex with `includegraphics`. Raster




Figure 3.1: Caption of the figure. Always cite the source of the image in the caption, unless you created the image yourself.

	Metric 1↑	Metric 2↑	Metric 3↓	Metric 4↓
Method 1	23.1388	0.8611	0.2479	0.1846
Method 2	28.3813	0.9353	0.1831	0.1047
Ours	26.2468	0.9257	0.1509	0.0915

Table 3.1: Caption of the Table

images should only be used in your thesis when an image is shown, that is naturally rendered at a given resolution, e.g., a result image of a rendering algorithm.

3.5 Tables

Table placement is best done with the default `table` environment. To reference a table use `tblref`, for example Table 3.1.

Chapter 4

Implementation

While the previous section aims to describe your method in a platform independent way, this section presents all the details to reproduce your work.

4.1 Listings

Not all theses require the inclusion of source code. If it is necessary, Code placement is best done with the `lstlisting` environment.

```
1 # Comment
2 def function(parameter):
3     print("This is an example Python listing")
4     return True
```

Listing 4.1: Caption of the Python Listing

```
1 // Comment
2 bool function(int parameter){
3     std::cout << "This is an example C++ listing" << std::endl;
4     return true;
5 }
```

Listing 4.2: Caption of the C++ Listing

To reference a listing use `lstref`, for example Listing 4.1.

Chapter 5

Results

Describe your qualitative and quantitative evaluations and present your results. Discuss the limitations of your work.

Chapter 6

Conclusions


Give a brief summary of your work and your contributions. Afterwards, briefly outline potential opportunities for future work.

6.1 Tips

As simple package to keep track of unfinished areas of the thesis is the `todonotes` package.

Inline TODO

This is
a simple
TODO
note



Create a list of remaining TODOs with `listoftodos`. Don't forget to remove the `listoftodos` before submitting your thesis.

Todo list

citation	1
This is a simple TODO note	11
Inline TODO	11
Figure: Missing figure	11

List of Figures

3.1 Caption of the Figure as listed in listoffigures 6

LIST OF FIGURES

List of Tables

3.1 Caption of the Table as listed in listoftables 6

LIST OF TABLES

List of Listings

4.1	Caption of the Python Listing	7
4.2	Caption of the C++ Listing	7

LIST OF LISTINGS

Bibliography

- [1] Juan Ballena-Ruiz, Juan Arcila-Diaz, and Victor Tuesta-Monteza. “Automated Detection and Counting of Gossypium barbadense Fruits in Peruvian Crops Using Convolutional Neural Networks”. In: *AgriEngineering* 7.5 (2025), p. 152. ISSN: 2624-7402. DOI: 10.3390/agriengineering7050152.
- [2] Nicolas Carion et al. *SAM 3: Segment Anything with Concepts*. 2026. arXiv: 2511.16719 [cs.CV]. URL: <https://arxiv.org/abs/2511.16719>.
- [3] G. Ramesh Chandra, K. Sri Rama Murthy, S. Archana, B. Sneetha, G. Nishanth, M. Navya, and R. Sai Charan. “Fruit Detection and Counting Using Deep Learning”. In: *2025 Eighth International Conference on Image Information Processing (ICIIP)*. 2025, pp. 175–182. DOI: 10.1109/ICIIP68302.2025.11346118.
- [4] Jordi Gené-Mola, Ricardo Sanz-Cortiella, Joan R. Rosell-Polo, Josep-Ramon Morros, Javier Ruiz-Hidalgo, Verónica Vilaplana, and Eduard Gregorio. “Fruit detection and 3D location using instance segmentation neural networks and structure-from-motion photogrammetry”. In: *Computers and Electronics in Agriculture* 169 (2020), p. 105165. ISSN: 0168-1699. DOI: 10.1016/j.compag.2019.105165.
- [5] Nicolai Häni, Pravakar Roy, and Volkan Isler. “A comparative study of fruit detection and counting methods for yield mapping in apple orchards”. In: *Journal of Field Robotics* 37.2 (Aug. 2019), pp. 263–282. ISSN: 1556-4967. DOI: 10.1002/rob.21902. URL: <http://dx.doi.org/10.1002/rob.21902>.
- [6] Xiangyu Hong, Linlin Qin, Chun Shi, and Gang Wu. “FruitGaussian: 3D Gaussian Splatting for Automated Fruit Counting in Natural Orchard”. In: *Proceedings of the 44th Chinese Control Conference (CCC)*. 2025, pp. 8004–8009. DOI: 10.23919/CCC64809.2025.11179457.
- [7] Fengze Li, Yangle Liu, Jieming Ma, Hai-Ning Liang, Yaochun Shen, Huangxiang Li, and Zhijing Wu. *CountingFruit: Language-Guided 3D Fruit Counting with Semantic Gaussian Splatting*. 2025. arXiv: 2506.01109 [cs.CV].
- [8] Xu Liu, Steven W. Chen, Shreyas Aditya, Nivedha Sivakumar, Sandeep Dcunha, Chao Qu, Camillo J. Taylor, Jnaneshwar Das, and Vijay Kumar. *Robust Fruit Counting: Combining Deep Learning, Tracking, and Structure from Motion*. 2018. arXiv: 1804.00307 [cs.CV]. URL: <https://arxiv.org/abs/1804.00307>.

BIBLIOGRAPHY

- [9] Xu Liu, Steven W. Chen, Chenhao Liu, Shreyas S. Shivakumar, Jnaneshwar Das, Camillo J. Taylor, James Underwood, and Vijay Kumar. “Monocular Camera Based Fruit Counting and Mapping With Semantic Data Association”. In: *IEEE Robotics and Automation Letters* 4.3 (2019), pp. 2296–2303. ISSN: 2377-3774. DOI: 10.1109/lra.2019.2901987. URL: <http://dx.doi.org/10.1109/LRA.2019.2901987>.
- [10] Lukas Meyer, Andreas Gilson, Ute Schmid, and Marc Stamminger. “FruitNeRF: A Unified Neural Radiance Field based Fruit Counting Framework”. In: *2024 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. 2024, pp. 1–8. DOI: 10.1109/IROS58592.2024.10802065.
- [11] Ben Mildenhall, Pratul P. Srinivasan, Matthew Tancik, Jonathan T. Barron, Ravi Ramamoorthi, and Ren Ng. “NeRF: representing scenes as neural radiance fields for view synthesis”. In: *Commun. ACM* 65.1 (Dec. 2021), pp. 99–106. ISSN: 0001-0782. DOI: 10.1145/3503250. URL: <https://doi.org/10.1145/3503250>.
- [12] Juan C. Miranda, Jordi Gené-Mola, Manuela Zude-Sasse, Nikos Tsoulas, Alexandre Escolà, Jaume Arnó, Joan R. Rosell-Polo, Ricardo Sanz-Cortiella, José A. Martínez-Casasnovas, and Eduard Gregorio. “Fruit sizing using AI: A review of methods and challenges”. In: *Postharvest Biology and Technology* 206 (2023), p. 112587. ISSN: 0925-5214. DOI: 10.1016/j.postharvbio.2023.112587.
- [13] Jianyuan Wang, Minghao Chen, Nikita Karaev, Andrea Vedaldi, Christian Rupprecht, and David Novotny. “VGGT: Visual Geometry Grounded Transformer”. In: *2025 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 2025, pp. 5294–5306. DOI: 10.1109/CVPR52734.2025.00499.

Declaration

I certify that I have written this thesis without outside help and without using sources other than those indicated, and that the thesis has not been submitted in the same or similar form to any other examination authority and has not been accepted by them as part of an examination. All statements that have been taken over verbatim or in spirit are marked as such.

During the preparation of this work, AI technologies were used to assist in the writing process. Specifically, <AI Tool 1 (Company, City, State, Country)> was used in order to check for grammar and style consistency, and <AI Tool 2 (Company, City, State, Country)> was used in order to assist with rephrasing and improving readability. After using these tools, the manuscript was carefully reviewed and the content was edited as needed. No tools or services were used for content generation. The content of this thesis was entirely created by myself, and I ensured that all material presented reflects my original work and research. I take full responsibility for the content of the thesis.

I agree that the thesis may be published and that it may be referred to in scientific publications.

The Friedrich-Alexander-Universität Erlangen-Nürnberg, represented by the Professorship for Cognitive Computer Vision, is granted a (non-exclusive) right to use this work and the programs created in connection with it.

Erlangen, 01.10.2026

(Alexander Link)